Multi-Population Concordance Study of the GenomeGuard Pharmacogenomic Engine Against PharmCAT Across 2,600 Diverse Genomes

Aditya Yadav (yadav.adi0718@gmail.com)[1], Aman Goyal (amangoyal5994@gmail.com)[1], Gaurav Yadav (gauravyadav6698@gmail.com)[1]

1. GenomeGuard Research Institute

Abstract

Background: Pharmacogenomic (PGx) testing holds transformative potential for precision medicine by preemptively matching patients with optimal drug therapies. However, existing genomic interpretation pipelines are often computationally intensive and rely heavily on complex containerized environments, posing significant barriers to population-scale implementation. We present a comprehensive benchmarking of GenomeGuard, a native Python PGx engine designed for extreme throughput, against the gold-standard PharmCAT reference implementation across 2,600 diverse genomes.

Methods: Star-allele diplotypes for 13 Tier-1 pharmacogenes were extracted from high-coverage (30×) whole-genome sequencing data representing five major global superpopulations from the 1000 Genomes Project. Concordance between GenomeGuard v2.0 and PharmCAT was rigorously assessed utilizing both exact-match agreement and multi-class Cohen's kappa (κ) to robustly account for the complex distributions of multi-class diplotype definitions across differing ancestries. Computational throughput and memory overhead were concurrently benchmarked.

Results: Across 33,800 independent gene-sample comparisons, GenomeGuard demonstrated 99.69% overall exact diplotype concordance with PharmCAT. In 11 of the 13 evaluated pharmacogenes, GenomeGuard achieved perfect concordance (κ = 1.000). A localized drop in exact-match concordance was uniquely observed for UGT1A1 within the African (AFR) superpopulation (90.15% concordance, κ=0.949). Algorithmic mismatch triage decisively traced this to underlying allele-definition divergence regarding the thresholding of TA-repeats (differentiating *28 vs *37 alleles) rather than algorithmic failure. Computationally, GenomeGuard achieved a 56× throughput speedup over the Dockerized PharmCAT pipeline (14.1 vs 0.25 samples/second) while utilizing minimal memory (<1MB peak).

Conclusions: GenomeGuard produces clinically equivalent diplotype calls to the PharmCAT tool, closely aligning with CPIC guidelines, while offering dramatically faster execution suitable for real-time clinical decision support and EHR integration across diverse global populations.

Keywords: *Pharmacogenomics, Concordance Study, GenomeGuard, PharmCAT, Precision Medicine, 1000 Genomes, Cohen's Kappa, Bioinformatics*

1. Introduction

Adverse drug reactions (ADRs) are a leading cause of morbidity and mortality worldwide, imposing substantial clinical and economic burdens on global healthcare systems [1]. Pharmacogenomics (PGx), the study of how an individual's genetic variations influence their response to pharmacological agents, offers a powerful mechanism to mitigate ADRs and optimize therapeutic efficacy [2]. By interrogating specific genetic loci—predominantly genes encoding cytochrome P450 (CYP) enzymes, drug transporters, and human leukocyte antigens (HLAs)—clinicians can preemptively identify patients at risk for severe toxicity or therapeutic failure. Organizations such as the Clinical Pharmacogenetics Implementation Consortium (CPIC) and the Dutch Pharmacogenetics Working Group (DPWG) provide highly curated, evidence-based clinical guidelines translating these genotype-to-phenotype relationships into actionable dosing recommendations [3].

In PGx, variations are canonically categorized into 'star alleles' (e.g., CYP2C19*2), which represent distinct functional haplotypes defined by the Pharmacogene Variation Consortium (PharmVar) [4]. The computational translation of raw next-generation sequencing (NGS) data into accurate star-allele diplotypes (and their associated metabolic phenotypes) remains a formidable bioinformatics challenge [5]. The Pharmacogenomics Clinical Annotation Tool (PharmCAT), developed jointly by PharmGKB and St. Jude Children's Research Hospital, is currently regarded as the gold-standard open-source reference implementation for this translation [6]. However, standard implementations are computationally intensive, often reliant on virtualization or containerized environments (like Docker) that introduce substantial runtime overhead, creating processing bottlenecks when deployed at population scale or integrated directly into low-latency clinical environments such as Electronic Health Records (EHR) [7].

In this study, we introduce and benchmark GenomeGuard v2.0, an optimized, cloud-native PGx engine designed for extreme throughput without sacrificing clinical fidelity. We present a massive-scale, multi-population concordance assessment of GenomeGuard against PharmCAT across 2,600 high-coverage (30×) whole genomes from the 1000 Genomes Project [8]. By assessing concordance across diverse global ancestries, we validate GenomeGuard's utility as a robust, highly scalable, and equitable solution for enterprise clinical precision medicine.

2. Methods

2.1 Study Dataset and Multi-Population Cohort

To ensure a rigorous and diverse evaluation, we utilized the New York Genome Center (NYGC) 30× high-coverage whole-genome sequencing (WGS) dataset derived from the expanded 1000 Genomes Project cohort [8]. We randomly selected 2,600 samples representing five major global superpopulations: African (AFR), Admixed American (AMR), East Asian (EAS), European (EUR), and South Asian (SAS). This robust demographic sampling is crucial, as PGx allele frequencies are known to vary significantly by biogeographical ancestry [9].

2.2 Pharmacogenomic Variant Panel and Preprocessing

Our analysis targeted distinct, defining variant positions across 13 Tier-1 pharmacogenes amenable to direct head-to-head comparison (ABCG2, CYP2B6, CYP2C19, CYP2C9, CYP3A4, CYP3A5, DPYD, IFNL3, NUDT15, SLCO1B1, TPMT, UGT1A1, and VKORC1). These specific loci were algorithmically selected to perfectly align with CPIC star-allele defining variants [3]. Three additional Tier-1 genes (CYP1A2, CYP2C8, NAT2) were assessed only via internal synthetic unit tests and excluded from external concordance statistics because PharmCAT's current allele-definition files do not produce callable results for these genes from short-read VCF data. Raw variant calls were extracted in GRCh38 coordinates using bcftools (v1.9) from the overarching NYGC multi-sample VCFs, producing individual, single-sample VCFs optimized for parallelized downstream ingestion.

2.3 Tool Benchmarking Framework

We directly compared GenomeGuard v2.0 against the latest iteration of PharmCAT. GenomeGuard operates as a native Python application, executing proprietary, high-speed VCF parsing and deterministic CPIC-aligned star-allele logic. Conversely, PharmCAT was executed via its official Docker container ('pgkb/pharmcat:latest'), utilizing its standard sequential pipeline: the Named Allele Matcher, followed by the Phenotyper, and concluding with the

Reporter [6]. All performance benchmarking was conducted on an isolated Windows subsystem environment to simulate a standard local clinical workstation.

2.4 Statistical and Concordance Analysis

Pairwise concordance was evaluated initially using exact-match agreement of canonical, sorted diplotypes. Because standard binary agreement metrics are insufficient for complex multi-class diplotype assignments, inter-rater reliability was rigorously quantified utilizing a multi-class Cohen's kappa (κ) [10]. The κ statistic was computed from the full N×N contingency matrix of all observed diplotype categories, effectively correcting for chance agreement across highly imbalanced allele distributions (the 'kappa paradox') [11]. Additionally, binary sensitivity, specificity, and F1 scores were computed per-gene, where 'positive' was defined as the detection of any non-reference (non-*1/*1) diplotype.

2.5 Algorithmic Mismatch Triage

To objectively categorize and interpret exact-match discordances, we implemented an automated, deterministic triage algorithm, ensuring symmetric classification. Mismatch categories were strictly defined as: (1) Missing Coverage (insufficient read depth), (2) Notation Divergence (equivalent biological alleles named using conflicting nomenclatures), (3) Allele Definition Versioning (divergence in the underlying structural definition of specific star alleles), and (4) Logic Differences (fundamental disagreements in phasing or combination logic).

3. Results

3.1 Diplotype Concordance and Non-Reference Detection

In total, 33,800 independent gene-sample comparisons were executed. GenomeGuard and PharmCAT achieved a highly robust overall exact diplotype concordance of 99.69%. Both tools demonstrated near-perfect harmony in detecting the presence of any non-reference diplotype, yielding a binary F1 score of 1.000 across all 13 evaluated pharmacogenes. When enforcing the stricter standard of exact multi-class diplotype matching, 11 of the 13 genes retained perfect concordance ($\kappa = 1.000$), reinforcing the high reliability of GenomeGuard's internal calling logic (Table 1).

Table 1. Comprehensive Per-Gene Concordance Metrics between GenomeGuard and PharmCAT.

Pharmacogene	Samples (N)	Exact Match (%)	Multi-Class κ	Binary F1	Mismatches
ABCG2	2600	100.00%	1.000	1.000	0
CYP2B6	2600	99.46%	0.991	1.000	14
CYP2C19	2600	100.00%	1.000	1.000	0
CYP2C9	2600	100.00%	1.000	1.000	0
CYP3A4	2600	100.00%	1.000	1.000	0
CYP3A5	2600	100.00%	1.000	1.000	0
DPYD	2600	100.00%	1.000	1.000	0
IFNL3	2600	100.00%	1.000	1.000	0
NUDT15	2600	100.00%	1.000	1.000	0
SLCO1B1	2600	100.00%	1.000	1.000	0
TPMT	2600	100.00%	1.000	1.000	0
UGT1A1	2600	96.50%	0.949	1.000	91
VKORC1	2600	100.00%	1.000	1.000	0

3.2 Population Equity and the UGT1A1 Divergence

Concordance remained exceedingly high across the AMR, EAS, EUR, and SAS superpopulations. However, statistical analysis revealed a highly localized drop in exact-match concordance specifically for the UGT1A1 gene within the African (AFR) superpopulation, falling to 90.15% (κ=0.949).

Application of the algorithmic mismatch triage revealed that of the 91 total UGT1A1 mismatches observed, 90 were singularly attributable to Allele Definition Versioning. This divergence centers on the highly polymorphic TA-repeat region in the UGT1A1 promoter [12]. GenomeGuard and PharmCAT employ differing categorical thresholds for distinguishing the *28 allele (7 TA repeats) from the *37 allele (8 TA repeats). Because the *37 allele is predominantly found in individuals of African descent, this definition skew manifested almost exclusively within the AFR cohort [9]. This highlights the importance of standardized allele definitions in multi-population PGx, rather than representing a fundamental flaw in algorithmic variant detection.

3.3 Computational Throughput and Scalability

Performance benchmarking revealed dramatic differences in computational efficiency. Operating natively, GenomeGuard processed the entire 2,600-sample cohort in just 184.6 seconds, achieving a sustained throughput of 14.1 samples per second. Furthermore, its lightweight architecture maintained a peak memory footprint of under 500 KB.

In stark contrast, PharmCAT averaged 0.25 samples per second. This 56× throughput advantage for GenomeGuard is heavily attributable to its targeted, pre-filtered position scanning and the avoidance of virtualization overhead. PharmCAT's requirement to initialize a Docker container and a Java Virtual Machine (JVM) for each sequential sample creates a significant processing bottleneck. While batch-mode execution of PharmCAT may mitigate some initialization overhead, GenomeGuard's native architecture inherently provides superior scalability for real-time, high-volume clinical applications.

4. Discussion

The successful integration of pharmacogenomics into routine clinical practice requires software infrastructure that is not only highly accurate and aligned with CPIC guidelines but also capable of scaling to population-level demands without imposing exorbitant computational costs [7]. This massive-scale concordance study robustly affirms that GenomeGuard reliably produces star-allele diplotype calls that mirror the established PharmCAT standard.

The minor discordances observed, particularly in UGT1A1, underscore a known challenge in the PGx community: the reliance on continuously evolving allele definitions [4]. Because multi-class κ is highly sensitive to the marginal distribution of prevalent alleles (the 'kappa paradox'), raw concordance rates and F1 scores remain critical for contextualizing agreement on rare, structurally complex variants like TA-repeats [11]. The ability of GenomeGuard to perfectly isolate and detect the presence of non-reference variants (F1 = 1.000) confirms its clinical safety baseline.

5. Conclusions

GenomeGuard v2.0 demonstrates robust analytical concordance against PharmCAT across globally diverse ancestries, confirming its viability as a reliable, CPIC-aligned PGx engine. Its high-velocity architecture, devoid of heavy containerization overhead, provides a highly scalable, equitable, and clinical-grade alternative to existing computational pipelines, strongly positioning it for real-time integration into precision medicine workflows.

6. References

[1] Pirmohamed, M., et al. (2004). Adverse drug reactions as cause of admission to hospital: prospective analysis of 18 820 patients. BMJ, 329(7456), 15-19.

[2] Relling, M. V., & Evans, W. E. (2015). Pharmacogenomics in the clinic. Nature, 526(7573), 343-350.

[3] Caudle, K. E., et al. (2014). Incorporation of pharmacogenomics into routine clinical practice: the Clinical Pharmacogenetics Implementation Consortium (CPIC) guideline development process. Current Drug Metabolism, 15(2), 209-217.

[4] Gaedigk, A., et al. (2018). The Pharmacogene Variation (PharmVar) Consortium: incorporation of the human cytochrome P450 (CYP) allele nomenclature database. Clinical Pharmacology & Therapeutics, 103(3), 399-401.

[5] McInnes, G., et al. (2021). Computational pharmacogenomics: challenges and opportunities. Clinical Pharmacology & Therapeutics, 109(4), 924-932.

[6] Sangkuhl, K., et al. (2020). PharmCAT: a pharmacogenomics clinical annotation tool. Clinical Pharmacology & Therapeutics, 107(1), 203-210.

[7] Volpi, S., et al. (2018). Research directions in the clinical implementation of pharmacogenomics: an overview of US programs and projects. Clinical Pharmacology & Therapeutics, 103(5), 778-786.

[8] Byrska-Bishop, M., et al. (2022). High-coverage whole-genome sequencing of the expanded 1000 Genomes Project cohort including 602 trios. Cell, 185(18), 3426-3440.

[9] Zhou, Y., et al. (2017). Population-scale distribution of pharmacogenomic variants. Clinical Pharmacology & Therapeutics, 102(4), 687-695.

[10] Cohen, J. (1968). Weighted kappa: Nominal scale agreement provision for scaled disagreement or partial credit. Psychological Bulletin, 70(4), 213-220.

[11] Feinstein, A. R., & Cicchetti, D. V. (1990). High agreement but low kappa: I. The problems of two paradoxes. Journal of Clinical Epidemiology, 43(6), 543-549.

[12] Innocenti, F., et al. (2014). All you need to know about UGT1A1 genetic testing for patients treated with irinotecan: a practitioner-friendly guide. Journal of Oncology Practice, 10(3), e122-e124.